

# Contents

|                                                                               |          |
|-------------------------------------------------------------------------------|----------|
| <b>Dialogue Data Generation and Evaluation: A Deep-Dive Synthesis</b>         | <b>1</b> |
| Table of Contents . . . . .                                                   | 1        |
| 1. Why This Problem Matters . . . . .                                         | 2        |
| 2. Historical Arc: Three Generations of Dialogue Data . . . . .               | 2        |
| 3. The Central Tension: Data Fidelity vs. Data Scale . . . . .                | 3        |
| 4. The Paradigm Flip: Synthetic Users Diverge From Real Ones . . . . .        | 4        |
| 5. Full Taxonomy of Dialogue Data Methods . . . . .                           | 4        |
| 6. Deep Dive: MultiWOZ and the Benchmark Era . . . . .                        | 6        |
| 7. Deep Dive: Soudani et al. 2024 — The Field’s Taxonomic Frame . . . . .     | 7        |
| 8. Deep Dive: Wang et al. 2025 — Synthetic Users Are Not Real Users . . . . . | 7        |
| 9. Deep Dive: Chen et al. 2024 — Diversity Matters, Predictably . . . . .     | 9        |
| 10. Supporting Papers: Quick Takes . . . . .                                  | 10       |
| 11. Open Problems and Research Frontiers . . . . .                            | 11       |
| 12. Synthesis: Where the Field Is Heading . . . . .                           | 12       |

## Dialogue Data Generation and Evaluation: A Deep-Dive Synthesis

**From Wizard-of-Oz to LLM-Driven Synthesis — measuring conversations at scale**

*Synthesized from UserLM.Data and UserLM.Evaluation Zotero collections (21 papers) March 2026 | Internal notes*

---

### Table of Contents

|                                                                                     |
|-------------------------------------------------------------------------------------|
| 1. <a href="#">Why This Problem Matters</a>                                         |
| 2. <a href="#">Historical Arc: Three Generations of Dialogue Data</a>               |
| 3. <a href="#">The Central Tension: Data Fidelity vs. Data Scale</a>                |
| 4. <a href="#">The Paradigm Flip: Synthetic Users Diverge From Real Ones</a>        |
| 5. <a href="#">Full Taxonomy of Dialogue Data Methods</a>                           |
| 6. <a href="#">Deep Dive: MultiWOZ and the Benchmark Era</a>                        |
| 7. <a href="#">Deep Dive: Soudani et al. 2024 — The Field’s Taxonomic Frame</a>     |
| 8. <a href="#">Deep Dive: Wang et al. 2025 — Synthetic Users Are Not Real Users</a> |
| 9. <a href="#">Deep Dive: Chen et al. 2024 — Diversity Matters, Predictably</a>     |
| 10. <a href="#">Supporting Papers: Quick Takes</a>                                  |
| 11. <a href="#">Open Problems and Research Frontiers</a>                            |
| 12. <a href="#">Synthesis: Where the Field Is Heading</a>                           |

---

## 1. Why This Problem Matters

Dialogue systems — particularly task-oriented dialogue (TOD) systems for customer service, bookings, and information retrieval — have become a critical industrial and research application. Yet they face a fundamental data problem: **genuine multi-turn task-oriented conversations are expensive to collect, sparse in the wild, and expensive to annotate**. A single human-human dialogue collection effort (like MultiWOZ’s Wizard-of-Oz setup) yields thousands of utterances but requires coordinating human annotators and users. Scaling beyond 10K dialogues via human effort becomes economically unfeasible.

This has spawned an entire research program around **dialogue data augmentation (DA)**: generating synthetic dialogues to supplement human-collected training data. The stakes are high. A 2-4 BLEU point improvement (typical DA gains) can translate to meaningful performance differences in downstream systems, and a botched augmentation strategy can introduce systematic biases that degrade performance on out-of-distribution data.

The field has evolved through three distinct technological paradigms, each with its own assumptions about *what constitutes good synthetic dialogue*. That evolution — and the recent discovery that synthetic users systematically diverge from real users — forms the core narrative of this synthesis.

---

## 2. Historical Arc: Three Generations of Dialogue Data

### Generation 1: Rule-Based and Wizard-of-Oz (2010-2019)

The foundational approach was human-in-the-loop: Wizard-of-Oz data collection, where annotators playing a “system” respond to real users. Budzianowski et al. (2020) operationalized this at scale with **MultiWOZ**, collecting 10,438 human-to-human dialogues across 7 domains (hotels, restaurants, attractions, trains, taxi, police, hospital) with 115K utterances and an average of 8 turns per dialogue. This became *the* benchmark.

Parallel to Wizard-of-Oz, rule-based augmentation emerged: template-based paraphrasing, back-translation (translating English to German and back, Soudani et al. 2023), and simple morphological variation. These methods yielded modest gains: 2-4 BLEU points. The implicit assumption was that *paraphrastics matter* — different surface realizations of the same intent improve robustness.

The critical limitation: rule-based methods don’t scale to new domains and struggle with phenomena like indirect requests, disfluencies, and pragmatic variation. Shalymov et al. (2017) exposed this with bAbI+, showing that MemN2N dialogue models trained on *perfectly clean* bAbI data plummeted from 100% to ~50% accuracy when tested on realistic phenomena like hesitations (21%), restarts (40%), and corrections (5%).

## Generation 2: Neural Augmentation (2018-2022)

The rise of sequence-to-sequence models and pointer-generator networks enabled more sophisticated augmentation. Seq2Seq models with copy attention, trained on large parallel corpora, could generate fluent paraphrases and back-translations. Pasali et al. (2022) applied this to disfluency generation (LARD: 95,992 synthetic disfluencies from TEDLIUM), improving ASR robustness by 2-3% WER while maintaining 81% human-rated naturalness.

Vlachos et al. (2024) provided the definitive empirical comparison: Easy Data Augmentation (EDA), back-translation, paraphrasing, and full dialogue-level generation on MultiWOZ. Results: - **Back-translation**: +2-3 BLEU - **Dialogue-level generation**: +3-5 BLEU - **Diminishing returns** beyond 2x augmentation ratio - **Task success improvements**: only 1-2 points (the critical gap)

Neural methods worked, but they revealed the evaluation crisis: BLEU, ROUGE, and BERTScore don't correlate with downstream task success. A model might gain 5 BLEU points from augmentation yet show only marginal improvement in real dialogue completion.

## Generation 3: LLM-Driven Synthesis (2023-present)

GPT-3, GPT-4, and similar LLMs made *in-context learning* feasible for dialogue generation. Soudani et al. (2024) surveys this transition: in-context learning with exemplars, fine-tuning with constraint-based prompts, and direct generation without task-specific training. LLM-based methods achieve 75–85% human preference over template-based approaches (Chavda & Bhattacharyya). PLACES (Chen et al. 2023) improves dialogue coherence by 12%. SynDG reaches 0.68 ROUGE-L.

But LLM methods introduce new failure modes: hallucination, inconsistent constraint adherence, and — most critically for UserLM — **systematic behavioral divergence from human users**. Wang et al. (2025) demonstrates this empirically: humans engage in depth-first questioning (explore one option fully), while LLM agents use breadth-first exploration (explore all options). This difference cascades into satisfaction ratings and dialogue structure.

---

## 3. The Central Tension: Data Fidelity vs. Data Scale

All three generations face a core trade-off: **fidelity vs. scale**.

Wizard-of-Oz dialogue is human-generated and pragmatically grounded, but slow and expensive. MultiWOZ's 10,438 dialogues took significant effort to collect. Going from 10K to 100K dialogues via human effort is economically infeasible.

Rule-based augmentation scales effortlessly (generate as many paraphrases as needed) but sacrifices fidelity. Template-based approaches produce stilted, repetitive variations.

Neural augmentation finds middle ground: generate thousands of fluent dialogues

in minutes, but with degraded task-alignment and inconsistent quality. Vlachos et al. (2024) showed that task success improvements plateau despite BLEU gains.

LLM-driven synthesis scales to arbitrary sizes and achieves high BLEU/ROUGE scores, but Wang et al. (2025) reveals the hidden cost: **LLM agents don't behave like humans**. They hallucinate differently, ask questions in different orders, and express uncertainty differently. Training on LLM-generated data may improve downstream NLG metrics while degrading dialogue system robustness in deployment.

The field's unspoken assumption — that synthetic data indistinguishable from human data *by surface metrics* will perform equivalently — is false. Fidelity isn't just linguistic fluency. It's pragmatic alignment.

---

#### 4. The Paradigm Flip: Synthetic Users Diverge From Real Ones

**Wang et al. (2025) is the turning point in this literature.** The paper directly compares 146 human participants against three LLM agents (GPT-4, Claude, Gemini) in 4 task scenarios, generating 1,856 human + 1,856 agent conversations. Findings:

- **Questioning strategy:** Humans use depth-first (fully explore 1 option), agents use breadth-first (explore all options). Statistically significant.
- **Error recovery:** Humans recover from errors iteratively; agents sometimes reinvent context mid-dialogue.
- **Hallucination awareness:** LLM agents show *lower* awareness of hallucinated information than humans.
- **Promise language:** Agents use higher frequency of commitment statements ("I will," "I can").
- **Cross-model variance:** Behaviors differ significantly across GPT-4, Claude, and Gemini implementations.

This has immediate implications for UserLM research: **if you train a dialogue system on LLM-generated user dialogues, you're training it on a user that doesn't exist**. The system learns to respond to breadth-first questioning and high-confidence utterances, then fails on real users who ask depth-first and express uncertainty.

Mannekote et al. (2025) partially addresses this. Their IndirectRequests dataset introduces pragmatic variation (indirect vs. direct requests) with linguistic criteria (appropriateness, unambiguity, world-understanding) and 0.67 inter-annotator agreement on naturalness. The key insight: **naturalness is about pragmatics, not surface paraphrase**. Their synthetic data improves task success by 3-4 points and reaches 0.67 inter-annotator agreement — modest but meaningful.

The implication: dialogue augmentation must be *behaviorally* informed, not just linguistically plausible.

---

#### 5. Full Taxonomy of Dialogue Data Methods

| Dimension                  | Approaches                          | Key Examples                               | Pros                                                  | Cons                                      |
|----------------------------|-------------------------------------|--------------------------------------------|-------------------------------------------------------|-------------------------------------------|
| <b>Generation Approach</b> | Rule-based templates                | Soudani et al. 2023                        | Fast, transparent, domain-agnostic                    | Stilted, low diversity                    |
|                            | Back-translation                    | Multiple papers                            | +2-3 BLEU, validated                                  | Limited scope, paraphrastics only         |
|                            | Seq2Seq + copy attention            | Passali et al. 2022 (LARD)                 | Fluent, disfluency-aware                              | Training overhead, hallucination          |
|                            | LLM in-context learning             | Soudani et al. 2024                        | High BLEU/ROUGE, zero-shot capable                    | Behavioral divergence (Wang et al.)       |
|                            | LLM fine-tuning (constraint-based)  | Soudani et al. 2024                        | Controllable, domain-aware                            | Requires labeled examples                 |
|                            | Token/word level                    | EDA (Vlachos et al. 2024)                  | Fine-grained control                                  | Limited fluency                           |
|                            | Sentence level                      | Paraphrasing (Cabezudo et al.)             | Preserves meaning                                     | Loss of dialogue coherence                |
|                            | Dialogue level                      | Dialogue-level generation (Vlachos 2024)   | +3-5 BLEU, pragmatic awareness possible               | Full retraining needed                    |
|                            | Domain-specific phenomena           | Disfluencies (Tian et al. 2021, LARD)      | ASR robustness (+2-3% WER)                            | Limited scope, labor-intensive annotation |
|                            | Stylized variation                  | KASDG (Li et al. 2023)                     | 0.62 BLEU improvement, human eval (4.2/5.0 coherence) | Requires explicit style labels            |
| <b>What Is Augmented</b>   | Code-switching (Winata et al. 2018) | Bilingual generation via pointer-generator | 55.19 BLEU, LM perplexity -10%                        | Language-pair specific                    |
|                            | BLEU / ROUGE                        | Ubiquitous                                 | Standardized, fast                                    | Doesn't correlate with task success       |
| <b>Evaluation Metric</b>   |                                     |                                            |                                                       |                                           |

| Dimension | Approaches                  | Key Examples                              | Pros                                     | Cons                                       |
|-----------|-----------------------------|-------------------------------------------|------------------------------------------|--------------------------------------------|
|           | BERTScore                   | Soudani et al. 2024 review                | Semantic alignment                       | Still misaligned with end-to-end metrics   |
|           | MAUVE                       | Soudani et al. 2024                       | Distribution matching                    | Requires reference set                     |
|           | Task success (JGA, slot F1) | Vlachos et al. 2024, Hudeček & Dusek 2023 | Directly relevant                        | Expensive, system-specific                 |
|           | Human preference            | Wang et al. 2025                          | Ground truth                             | High-cost, scale-limited (1,856 dialogues) |
|           | LLM cluster-agent diversity | Chen et al. 2024                          | Correlates with pre-training performance | Not yet validated on downstream NLG tasks  |

## 6. Deep Dive: MultiWOZ and the Benchmark Era

**Budzianowski et al. (2020)** published MultiWOZ, the canonical task-oriented dialogue dataset: 10,438 human-to-human dialogues collected via Wizard-of-Oz, 7 domains (hotels, restaurants, attractions, trains, taxi, police, hospital), 29 domain-specific slots, 115K utterances, average 8 turns per dialogue.

MultiWOZ became the benchmark *because* it was expensive to collect. It forced the field to focus: instead of proposing new datasets, papers propose methods that improve performance *on MultiWOZ*. The reproducibility benefit is immense. But it also ossified the field around a specific notion of “task-oriented dialogue” — formal, goal-driven, slot-filling-centric.

**Eric et al. (2020)** followed with MultiWOZ 2.1, fixing a critical flaw: ~40% of belief states in the original dataset were corrupted, leading to cascading errors in downstream state tracking. Eric et al. cleaned 4,183 dialogues, added dialogue act tags, and improved the state tracking baseline from 67% to 71% joint goal accuracy (JGA). BERTScore improved by 2-3 points. This is essential quality control — the cleaned version became the de facto gold standard.

The gap between MultiWOZ 1.0 and 2.1 reveals a hard lesson: **annotation quality, not just quantity, drives downstream performance**. It also shows why benchmark curation is expensive. One well-curated dataset worth \$100K of labor beats ten hastily annotated datasets worth \$10K each.

For UserLM specifically, MultiWOZ’s canonical status creates an evaluation bottleneck. Most papers report results on MultiWOZ’s *system* side (response generation, state tracking). Few probe the *user* side (what makes a realistic user utterance?). Mannekote et al. (2025) is an exception, introducing pragmatic variation. Wang et al. (2025) sidesteps MultiWOZ entirely, collecting new data to study human vs. agent behavior. This divergence — benchmark-centric papers vs. behavior-focused papers — hints at a field in transition.

---

## 7. Deep Dive: Soudani et al. 2024 — The Field’s Taxonomic Frame

**Soudani et al. (2024)** survey covers 100+ papers on conversational data generation, providing the definitive taxonomic frame for the field. Their categorization:

1. **Template-based** (rule-based paraphrasing, slot-filling)
2. **Rule-based neural** (pattern matching + neural reranking)
3. **Pre-LLM neural** (Seq2Seq, pointer-generator networks)
4. **LLM-driven** (in-context learning, fine-tuning with constraint-based prompts)

Each category trades off **controllability** (templates high, LLMs low) against **fluency** (templates low, LLMs high).

The evaluation section is the survey’s load-bearing component. Soudani et al. review metrics across four categories:

- **Surface-level:** BLEU, ROUGE, METEOR (penalize paraphrases, outdated)
- **Semantic:** BERTScore, chrF (contextualized embeddings, better but still not task-aligned)
- **Distribution:** MAUVE (matches distribution shape but requires reference set)
- **Task-specific:** BLEU/ROUGE for generation, JGA/slot F1 for state tracking

Their damning conclusion: **“No single metric reliably predicts downstream task success.”** This is the field’s central crisis. You can have a 5-point BLEU improvement with a 1-point JGA decrease. The metrics are orthogonal.

Soudani et al. propose multi-faceted evaluation: combine semantic (BERTScore), distributional (MAUVE), and task-specific (JGA) metrics. But this increases evaluation cost. For practitioners, the immediate implication is grim: **you must evaluate on your target task.** Generic metrics can’t proxy.

Domain-specific synthetic data challenges are mentioned but underdeveloped. Financial dialogues need different augmentation than restaurant bookings. Soudani et al. note this but don’t resolve it, pointing to a gap between the survey’s breadth and its depth on application.

---

## 8. Deep Dive: Wang et al. 2025 — Synthetic Users Are Not Real Users

**Wang et al. (2025)** directly compares human and LLM agent behavior in task-oriented conversations, a first-of-its-kind study with immediate implications for

UserLM.

**Setup:** 146 human participants + 3 LLM agents (GPT-4, Claude, Gemini)  $\times$  4 task scenarios (hotel booking, restaurant reservation, flight search, tech support) = 3,712 total dialogues (1,856 human, 1,856 agent). Metrics: dialogue success, user satisfaction, reasoning patterns.

**Key findings:**

1. **Questioning strategy** (statistically significant):
  - Humans: depth-first. Fully explore one option (e.g., all hotels in downtown) before moving to alternatives.
  - Agents: breadth-first. Sample across all options (one downtown, one midtown, one airport).
  - Implication: training on agent-generated dialogues teaches your system *the wrong exploration order*.
2. **Error recovery:**
  - Humans iterate: “That price is too high. What about next week?”
  - Agents sometimes reinvent: “Let me show you a different property entirely.”
  - Implication: system trained on agent data may not learn to refine within constraints.
3. **Hallucination awareness:**
  - Humans catch hallucinations mid-dialogue: “Wait, you said it had Wi-Fi but I don’t see that listed.”
  - Agents show lower awareness, sometimes doubling down on hallucinated claims.
  - Implication: agent-trained systems may propagate errors rather than acknowledge limits.
4. **Promise language:**
  - Agents: 15% higher frequency of commitment statements (“I will,” “I can,” “I’ll get you”).
  - Humans: more tentative (“I think,” “It looks like”).
  - Implication: agent-trained systems may over-promise.
5. **Cross-model variance:**
  - GPT-4, Claude, and Gemini show significantly different behavior profiles.
  - No single “LLM user” — each model is a different synthetic user.
  - Implication: results don’t generalize across model families.

**Methodological strength:** 1,856 human dialogues is substantial enough to establish statistical significance. Human satisfaction ratings (7.8/10 for humans vs. 6.4/10 for agents on average) directly measure the downstream cost.

**Limitation:** The paper doesn’t propose a solution, only diagnoses the problem. It doesn’t explore whether *mixed training* (human + agent data) could mitigate the bias, or whether behavioral fine-tuning of agents would help.

**Impact for UserLM:** This paper makes clear that UserLM systems trained on synthetic users will generalize poorly to real users *unless behavioral alignment is explicitly engineered*. The field’s implicit assumption — that fluency and task success correlate — is violated.



---

## 9. Deep Dive: Chen et al. 2024 — Diversity Matters, Predictably

**Chen et al. (2024)** tackles the diversity question: does the diversity of synthetic pre-training data correlate with downstream performance? And if so, how?

**Method:** They define **LLM cluster-agent diversity** as follows: - Cluster 620K Wikipedia topics using topic models. - Sample topics to create synthetic pre-training datasets with varying diversity. - Train 350M and 1.4B parameter models on synthetic data with controlled diversity. - Measure downstream performance on 8 benchmarks.

### Key findings:

1. **Diversity metric correlates with performance:**
  - LLM cluster score (higher = more diverse) positively correlates with pre-training perplexity and downstream task performance.
  - Effect size is modest but consistent across model scales.
2. **Optimal synthetic:real ratio is ~30:70** for supervised fine-tuning (SFT):
  - 100% synthetic: poor generalization.
  - 70% synthetic, 30% real: best downstream performance.
  - 30% synthetic, 70% real: diminishing gains from synthetic data.
3. **Diversity matters more for pre-training than fine-tuning** in smaller models:
  - 350M parameter model shows strong diversity sensitivity.
  - 1.4B parameter model shows weaker sensitivity (possibly due to over-parametrization absorbing noise).
4. **Scale effects are nonlinear:**
  - Small models need diversity; large models can afford some homogeneous data.
  - This hints at a capacity-dependent phenomenon.

**Measurement caveat:** The diversity metric (topic coverage) is task-agnostic. It's unclear whether topic diversity translates to dialogue-specific diversity (e.g., diverse reasoning patterns, pragmatic variation, or error types). The paper doesn't validate this transfer.

**Impact for UserLM:** Chen et al. suggest that **synthetic user diversity matters for downstream performance**. A UserLM trained on 1,000 diverse user personas will generalize better than one trained on 1,000 variations of the same user. But the paper doesn't specify what "diverse user personas" means operationally. Is it diverse goals? Diverse communication styles? Diverse error patterns? Wang et al. (2025) hints at the answer (behavioral diversity), but Chen et al. doesn't explicitly validate this.

The 30:70 ratio is practical guidance. If you're building a UserLM, aim for ~30% synthetic, 70% real human dialogue. But this assumes synthetic data quality, which Wang et al. casts doubt on.

## 10. Supporting Papers: Quick Takes

**Mannekote et al. (2025):** Introduces pragmatic variation (indirect requests) with linguistic criteria. Task success +3-4 points. Inter-annotator agreement on naturalness 0.67. Key insight: naturalness is pragmatics, not paraphrase. This bridges Soudani’s taxonomic frame and Wang’s behavioral findings.

**Vlachos et al. (2024):** Empirically compares EDA, back-translation, paraphrasing, dialogue-level generation. Dialogue-level generation wins (+3-5 BLEU) but task success gains modest (1-2 points). Diminishing returns at 2x augmentation. Essential for practitioners.

**Tian et al. (2021) — TOD-DA:** Two-stage augmentation (enrichment + speech quality simulation). Improves spoken robustness by 5-8% F1 while preserving written performance. Domain-specific (ASR channel).

**Passali et al. (2022) — LARD:** 95,992 synthetic disfluencies via Seq2Seq. Improves ASR by 2-3% WER, 81% human naturalness. Disfluency-focused, not general dialogue.

**Li et al. (2023) — KASDG:** Stylized dialogue generation with feature-guided selection + contrastive learning. 0.62 BLEU improvement, 4.2/5.0 coherence and 4.1/5.0 style appropriateness. Shows explicit style guidance works.

**Winata et al. (2018):** Code-switching via pointer-generator. 55.19 BLEU on bilingual data, LM perplexity -10%. Multilingual-specific.

**Cabezudo et al.:** Paraphrase generation for AMT-to-text (Portuguese). T5-based, +3.1 BLEU, diminishing returns beyond 8 paraphrases/instance. Validates paraphrase filtering.

**Shorten et al. (2021):** General text DA survey. Token-level vs. sentence vs. document-level. Back-translation and paraphrasing achieve 2-5% accuracy improvements. Dialogue is one domain among many.

**Andreas (2020) — GECA:** Fragment-based compositional augmentation for semantic parsing. 3-5% improvement, 10x faster than neural, 95%+ grammaticality. Not dialogue-specific, but shows promise for structured tasks.

**Shalymov et al. (2017):** bAbI+ dataset with realistic phenomena (hesitations, restarts, corrections). MemN2N drops to ~50% accuracy vs. 100% on clean bAbI. Establishes the clean-data problem that augmentation tries to solve.

**Hudeček & Dusek (2023):** LLM-based TOD pipeline (zero-shot JGA 51.6%, few-shot 54.8%). LLMs excel at generation, struggle with state tracking. Establishes LLM baseline for TOD.

**Zheng et al. (2023):** MT-Bench and Chatbot Arena. GPT-4 judge achieves >80% agreement with humans on open-ended questions. Identified position bias, verbosity bias, self-enhancement bias. Enables scalable dialogue evaluation but introduces new biases.

**Chavda & Bhattacharyya:** Covers PLACES (Chen et al. 2023) and SynDG. LLM methods achieve 75-85% human preference over templates. PLACES improves co-

herence by 12%, SynDG reaches 0.68 ROUGE-L. Bridges template-era to LLM-driven era.

---

## 11. Open Problems and Research Frontiers

**1. Behavioral alignment without data cost:** Wang et al. (2025) exposes behavioral divergence (breadth-first vs. depth-first, hallucination awareness, promise language). Solutions exist only in sketch form. Fine-tuning LLMs to match human behavior requires either labeled behavioral examples (expensive) or reinforcement learning from dialogue outcomes (unstable). Mannekote et al. (2025) shows that pragmatic variation helps, but it’s manually crafted, not learned. The challenge: how to make synthetic users behaviorally aligned without 10K human dialogues?

**2. Domain-specific diversity:** Chen et al. (2024) shows diversity matters, but their metric (topic coverage) is task-agnostic. What does diversity mean for a financial dialogue system vs. a tech support system? Is it diversity of financial products, or diversity of error types, or diversity of user sophistication? Soudani et al. (2024) notes this gap but doesn’t resolve it. Domain adaptation for dialogue augmentation remains an unsolved problem.

**3. Evaluation metric alignment:** Soudani et al. (2024) identifies the central crisis: BLEU, ROUGE, and BERTScore don’t correlate with task success. Proposed solutions (multi-faceted evaluation combining MAUVE, BERTScore, and JGA) require expensive downstream evaluation. Is there a middle ground — a metric that’s both cheap and predictive?

**4. Synthetic data decay:** As models are trained on synthetic data, quality degrades (a feedback loop known as “model collapse”). Soudani et al. alludes to this but doesn’t explore it. If you train an LLM on 70% synthetic data, then use that LLM to generate more synthetic data, how many generations can you sustain before quality collapses?

**5. Real vs. synthetic mixing strategies:** Chen et al. (2024) proposes 30:70 synthetic:real ratio, but this is empirically derived, not theoretically grounded. Does it hold across domains? Across model scales? Can you mix intelligently (e.g., synthetic for rare phenomena, real for common phenomena) or is simple ratios optimal?

**6. Cross-modal robustness:** Tian et al. (2021) and Passali et al. (2022) focus on spoken dialogue. Most augmentation work is text-only. How should augmentation strategies change for speech-to-text pipelines? Is the diversity metric different?

**7. User modeling at scale:** Wang et al. (2025) identifies specific behavioral differences. How do these cascade to system design? Should dialogue systems have user models that adapt to breadth-first vs. depth-first questioners? Should they learn to be more tentative to match human language? This is a systems problem, not just a data problem.

---

## 12. Synthesis: Where the Field Is Heading

The evidence points to a coherent synthesis, but it contradicts the field’s current direction.

**The field’s consensus position (circa 2024):** LLM-driven synthetic dialogue generation has “solved” data augmentation. Scale up the data, ensure linguistic fluency, and downstream performance follows. This reflects the success of large language models in NLP more broadly — scale and fluency beat everything.

**The actual evidence:** Scale without behavioral alignment is a productivity trap. Wang et al. (2025) shows that LLM-generated users don’t behave like real users. Breadth-first vs. depth-first. Hallucination-unaware vs. hallucination-catching. High promise language vs. tentative speech. These differences aren’t surface phenomena that disappear at scale; they’re systematic biases that propagate into deployed systems.

Mannekote et al. (2025) and Chen et al. (2024) together point to the solution: **behavioral diversity, not just linguistic diversity**. You need synthetic dialogue that’s pragmatically varied (indirect vs. direct requests), error-aware (catches hallucinations), and appropriately uncertain (tentative language). This is harder to engineer than fluent text, but it’s necessary.

The field’s next phase should be:

1. **Shift from BLEU to behavioral metrics:** Stop optimizing for surface-level fluency. Start measuring behavioral alignment (questioning strategy, error recovery patterns, confidence calibration). This requires human annotation but is necessary for evaluation.
2. **Pragmatically-aware generation:** Build augmentation methods that generate diverse user behaviors, not just diverse surface forms. This might mean prompting LLMs with behavioral exemplars, or training smaller models to perturb behavioral aspects of user utterances. Mannekote et al.’s pragmatic criteria provide a template.
3. **Diagnostic testing:** Adopt Wang et al.’s methodology. Test your synthetic user data against real users on controlled tasks. Measure divergence explicitly. Don’t assume fluency implies fidelity.
4. **Targeted mixing:** Rather than blanket 30:70 synthetic:real ratios, mix strategically. Synthetic data for rare phenomena (disfluencies, indirect requests), real data for common phenomena. This requires understanding your distribution, but it’s more efficient than uniform mixing.
5. **Behavioral fine-tuning:** Once you’ve identified behavioral divergences (Wang et al. study), fine-tune synthetic data generation to match. This might involve RL from human preference signals or constraint-based prompting. The cost is high, but it’s an investment in data quality.

**The unified theory:** Good synthetic dialogue is **diverse in behavior, not just in surface form**. Diversity (Chen et al.) matters because behavioral diversity improves

generalization. Wang et al. shows that LLMs default to a narrow behavioral space (breadth-first, high confidence). You must actively engineer behavioral variation.

For UserLM specifically, this means: synthetic user data should capture the full spectrum of human questioning strategies, error-handling behaviors, and linguistic confidence levels. A UserLM trained on such data will generalize to real users. One trained on LLM-generated users will fail on the first depth-first human.

**Timeline:** By 2026, the papers cited here should crystallize into a best-practice framework. Practitioners will stop reporting BLEU improvements as markers of success. Behavioral alignment will become the unit of assessment. The field will shift from “generate more diverse text” to “generate more diverse *behaviors*.” This is a subtle shift, but it’s foundational.

The hard part: behavioral diversity is expensive to annotate and harder to automate. But it’s necessary. The empirical evidence is now clear.

---

*End of synthesis. Document compiled from 21 papers in UserLM.Data and UserLM.Evaluation collections.*