

Contents

User Simulators & User Language Models: A Deep-Dive Synthesis	1
Table of Contents	1
1. Why This Problem Matters	2
2. Historical Arc: Four Generations of User Simulation	2
3. The Central Tension: State Alignment vs. Response Imitation	4
4. The Paradigm Flip: Flipping the Dialogue	6
5. Full Taxonomy of Current Approaches	7
6. Deep Dive: HumanLM (Wu et al., Stanford/NYU/Accenture, Feb 2026) . .	8
7. Deep Dive: Flipping the Dialogue / UserLM (Naous, Laban, Xu, Neville — Microsoft Research, Oct 2025)	10
8. Deep Dive: Know You First and Be You Better / USP (Wang et al., CUHK/NUS, ACL 2025)	13
9. Deep Dive: ChatBench (Chang, Anderson, Hofman — Microsoft Research, ACL 2025)	16
10. Jessie Lin's Framework: What Does It Take?	17
11. Supporting Papers: Quick Takes	19
12. Open Problems and Research Frontiers	21
13. Synthesis: Where the Field Is Heading	24
Appendix: Quick Reference Tables	26

User Simulators & User Language Models: A Deep-Dive Synthesis

From Response Imitation to State-Aligned Human Modeling

*Synthesized from the UserLM.Systems subcollection — Avishek's Zotero library
March 2026 | Internal notes, QED-Nano style*

Table of Contents

1. Why This Problem Matters: The RL Feedback Loop
2. Historical Arc: Four Generations of User Simulation
3. The Central Tension: State Alignment vs. Response Imitation
4. The Paradigm Flip: Flipping the Dialogue
5. Full Taxonomy of Current Approaches
6. Deep Dive: HumanLM (Wu et al., 2026)
7. Deep Dive: Flipping the Dialogue / UserLM (Naous et al., 2025)
8. Deep Dive: Know You First and Be You Better / USP (Wang et al., 2025)
9. Deep Dive: ChatBench (Chang et al., ACL 2025)
10. Jessie Lin's Framework: What Does It Take?
11. Supporting Papers: Quick Takes
12. Open Problems and Research Frontiers
13. Synthesis: Where the Field Is Heading

1. Why This Problem Matters

The fundamental bottleneck in training conversational AI systems is **feedback**. Reinforcement learning from human feedback (RLHF) requires humans in the loop — expensive, slow, and impossible to scale to the diversity of conversational scenarios a deployed system will encounter. The natural solution is a **user simulator**: a computational model that can play the human side of a conversation, generating realistic queries, follow-ups, clarifications, and reactions.

This problem is not new. But the LLM era has sharpened it considerably:

The RL environment problem. Before you can train a dialogue agent via RL, you need an environment to train *against*. Humans are the real environment, but they're not available at training time. User simulators are the proxy. Their quality directly caps the ceiling of what RL-trained dialogue systems can achieve — a good simulator enables better training signal; a bad simulator leads to policy collapse or artifacts that don't transfer to real users.

The evaluation problem. Static benchmarks (MMLU, MT-Bench, etc.) don't capture how AI systems perform when the human is actively shaping the conversation. A user who can't formulate their question well, who changes their mind mid-conversation, or who has implicit assumptions that the AI needs to surface — these dynamics are invisible to traditional evaluation. ChatBench (§9) makes this precise: AI systems that score identically on static benchmarks can differ enormously in user-AI collaborative accuracy.

The personalization problem. AI systems increasingly need to adapt to individual users — their style, expertise level, beliefs, preferences. Training and evaluating personalized systems requires a simulator that can *be* a specific person, not just a generic "user type." The USP work (§8) is the fullest attempt to date at this.

The compounding cost problem. Getting real human responses for diverse conversation trees (branching, multi-turn, counter-factual continuations) is combinatorially expensive. Synthetic user turns at scale are a prerequisite for many data augmentation and red-teaming workflows.

All of this converges on one question: **What does it mean to build a good user simulator, and how do you measure it?**

2. Historical Arc: Four Generations of User Simulation

Generation 1: Agenda-Based Simulators (pre-2020)

Hand-coded dialogue managers with explicit user *agendas* — structured representations of goals and sub-goals (e.g., book a flight from NYC to LAX, prefer morning departure). The simulated user works through the agenda item by item, generating surface-level natural language via templates or small seq2seq models.

- **Strengths:** Fully controllable; goal fidelity is near-perfect; interpretable.
- **Weaknesses:** Brittle natural language generation; no personality, emotion, or belief dynamics; cannot model the *unexpected* — topic drift, confusion, dissatisfaction.
- **Canonical examples:** PyDial, ConvLab, MultiWOZ-based simulators.

Generation 2: Neural Simulators (2020-2022)

End-to-end trained seq2seq models (GPT-2, T5, BART) trained on human conversation datasets to generate user turns given dialogue history. No explicit goal representation — the model implicitly learns response patterns.

- **Strengths:** More naturalistic language; can generalize to novel phrasings.
- **Weaknesses:** No explicit goal representation → poor goal fidelity; responses are statistically average rather than persona-specific; no coherent internal state across long conversations.

Generation 3: LLM-Prompted Simulators (2022-2024)

Instruction-tuned LLMs (GPT-3.5, GPT-4, Claude) prompted with a persona description and dialogue history. Zero-shot or few-shot simulation.

- **Strengths:** Flexible; quick to deploy; decent naturalness.
- **Weaknesses:** Jessy Lin (§10) catalogs five specific failure modes endemic to this generation:
 1. **Too knowledgeable** — the model’s vast training knowledge leaks through, creating “users” who know things their persona would not.
 2. **Perfect memory** — real users forget, lose context, and refer back imprecisely; prompted LLMs maintain perfect recall.
 3. **Too helpful** — the LLM’s instruction-following alignment makes it cooperate with the AI agent in ways that real users (who may be confused, adversarial, or just unhelpful) would not.
 4. **Inconsistent beliefs** — with no persistent internal state, the same “user” may hold contradictory views across turns or across sessions.
 5. **Shallow acting** — the model plays a *role* rather than *being* the person; the persona is a surface veneer, not a generative model of internal states.

Generation 4: Fine-Tuned User Language Models (2024-present)

Models trained *specifically* to simulate users — not assistants. The distinguishing architectural/training choices:

- **Base model, not instruction-tuned checkpoint** — instruction tuning bakes in assistant-like behavior (helpfulness, safety, formatting); user simulation requires *unlearning* these properties. Both UserLM (§7) and HumanLM (§6) make this case explicitly.
- **User-side data only** — trained on collections where only the *human* turns from real conversations are the training signal.
- **State-aligned vs. imitation** — the key methodological fork (§3).

This is the active research frontier. The three main technical papers in this collection (HumanLM, UserLM/Flipping, USP/Know You First) are all Generation 4 systems, each representing a different architectural bet.

3. The Central Tension: State Alignment vs. Response Imitation

This is the conceptual core of the field’s current argument, and **HumanLM** makes it most sharply.

Response Imitation (the naive approach)

Train the model to predict the next user utterance given conversation history. The training signal is: *did the model produce a response that looks like what a real user would say?*

- Loss function: cross-entropy on user turn tokens.
- Data: (context, user response) pairs.
- What the model learns: surface-level response distributions conditioned on conversational context.

The fundamental problem: Response imitation cannot distinguish between a user who says “What’s the price?” because they genuinely don’t know the price (low knowledge state) vs. a user who asks because they’re testing the system (high knowledge, adversarial stance). The words are identical; the internal state is completely different. A model trained only on responses will generate the right words but cannot model *why* a user produces those words — which means it will fail to generalize to novel conversational configurations.

Jessy Lin’s “shallow acting” failure mode is precisely this: the model mimics the surface of user behavior without having an underlying model that *generates* that behavior.

State Alignment (HumanLM’s proposal)

Train the model to maintain a coherent **internal state representation** that is consistent with what a real user in that situation would have, and *then* generate responses from that state.

The latent state space HumanLM proposes has **six psychologically grounded dimensions**:

Dimension	Description	Example
Belief	What the user thinks is true about the world	“I think the medication causes drowsiness”
Goal	What the user wants to achieve	“Get a specific recommendation, not general advice”

Dimension	Description	Example
Value	What the user cares about	"Privacy > convenience"
Stance	User's attitude toward the topic/AI	"Skeptical of AI-generated medical advice"
Emotion	Affective state	"Frustrated because previous AI response was vague"
Communication Style	How the user prefers to express themselves	"Terse, technical, dislikes verbose responses"

These six dimensions are not arbitrary — they map onto established psychological models of conversational behavior (Theory of Mind, speech act theory, etc.). The state is **latent** (not explicitly annotated in training data) and must be inferred.

Training signal: GRPO (Group Relative Policy Optimization), an RL objective. Rather than predicting the next token, the model is rewarded when: 1. The simulated user's internal state is consistent with what a human annotator would infer from context. 2. The generated response is consistent with the model's own stated internal state.

This forces state coherence across turns — the model cannot change its "goal" halfway through a conversation without a reason, because the reward signal penalizes state inconsistency.

Why This Matters Practically

The state alignment vs. imitation distinction has direct downstream consequences:

- **RL training quality:** An RL agent trained against a state-aligned simulator will face users who update their beliefs based on what the AI tells them (belief dynamics), who get frustrated when their goal isn't met (emotion dynamics), and who may not give the AI a second chance after a bad response (stance hardening). An imitation-based simulator lacks all of this; the RL-trained agent learns to produce good *local* responses but misses the long-horizon structure of user psychology.
- **Generalization:** State alignment provides the compositional structure needed to generalize to novel (belief, goal, value, stance, emotion, style) combinations. Imitation generalizes only to surface patterns seen in training.
- **Interpretability:** A state-aligned model can explain *why* it generated a particular response ("I asked for clarification because my goal state has high uncertainty about X"). Imitation models cannot.

The remaining question — which the field has not yet answered — is whether state alignment is *sufficient* or whether it also requires the right inductive biases in the

model architecture. GRPO provides the training framework; the representation learning question is open.

4. The Paradigm Flip: Flipping the Dialogue

Microsoft Research’s **Naous et al. (2025)** reframes the problem with a different conceptual emphasis. Rather than “how do we make user simulation better,” they ask: **“What if we built a language model whose primary role is the user side of dialogue — a purpose-built User LM?”**

This is a naming and framing choice that carries real technical content:

The Framing

Current LLM development pipelines optimize aggressively for *assistant* performance:

- Pre-training: mostly document-completion, not conversational.
- Instruction tuning: explicit fine-tuning to be helpful, harmless, honest.
- RLHF: further reinforcement of assistant-aligned behavior.

Every step bakes in assistant priors. Using such a model as a user simulator means you’re fighting against the model’s trained disposition at every step. The natural language of an over-helpful AI will never approximate the natural language of a confused, impatient, or curious human.

The Flipping paper’s prescription: **start from a base model checkpoint before instruction tuning** and fine-tune only on user-side conversational data. This preserves the model’s full language modeling capacity while avoiding the contamination of assistant alignment.

UserLM-8b

Their concrete artifact is **UserLM-8b**, built on top of Llama 3.1-8B (base checkpoint):

- Trained on user turns from **WildChat** — a large corpus of real human-AI conversations scraped from public deployments.
- **Not** trained on assistant turns.
- Evaluated on 6 intrinsic dimensions (see §7 for full details).

Key empirical finding: **base model > instruct model for user simulation across all 6 evaluation dimensions**. This is a strong, clean result that validates the paradigm flip framing — it’s not just a theoretical argument.

The Broader Claim

The paper positions User LMs as a **new model class** alongside:

- Chat/assistant LLMs (optimized for AI side of dialogue)
- Embedding models (optimized for representation)
- Code LLMs (domain-specialized)

The argument is that the AI ecosystem currently has robust infrastructure for the AI side of conversations, and almost nothing for the human side — despite the fact that the human side is what drives downstream system quality. User LMs fill this gap.

This framing has implications for **how we think about the AI development pipeline**: training a better AI system might require first training a better *user* system to evaluate and train against.

5. Full Taxonomy of Current Approaches

The field can be organized along several axes:

Axis 1: What the simulator is conditioned on

Type	Conditioning	Papers
Persona-conditioned	Explicit profile description	USP (Wang et al.)
Implicit-profile	Profile extracted from conversation history	USP, HumanLM
State-conditioned	Latent internal state	HumanLM
Unconditioned	Only dialogue history	UserLM/Flipping, ChatBench
Goal-conditioned	Task/information goal	ChatBench, agenda-based s

Axis 2: Training approach

Method	Description	Papers
SFT on user turns	Standard supervised fine-tuning on human side of conversations	UserLM/Flipping
SFT + RL	SFT initialization + reward-model-based RL	HumanLM (GRPO)
Conditional SFT	SFT conditioned on extracted profile	USP
RLCC	RL with Cycle Consistency — profile → dialogue → profile must round-trip	USP
In-context learning	Prompted inference, no fine-tuning	Lin’s Gen-3 critique

Axis 3: What capability is being optimized

Capability	Description
Distributional fidelity	Does the simulator’s language distribution match real users? (PPL)
Goal fidelity	Does the simulated user actually pursue their stated goal?
Profile consistency	Do utterances reflect the user’s persona accurately? (DPC)
State coherence	Are internal states consistent across turns?

Capability	Description
Belief updating	Does the user update beliefs when given new information?
Emotional realism	Are emotional responses appropriate to context?

Axis 4: Evaluation methodology

Method	Papers
Perplexity (PPL) on held-out human turns	UserLM/Flipping
DPC (Dialogue Profile Consistency) via FactScore	USP
ROUGE/BLEU similarity to reference responses	Multiple
Human evaluation	Multiple
LM-judge evaluation (GPT-4)	Multiple
Downstream task correlation	ChatBench

Axis 5: Scope of modeling

Scope	Description	Papers
Single-turn	Quality of individual user turns	Most early work
Multi-turn coherence	Consistency across a dialogue session	HumanLM, USP
Cross-session identity	Same “user” behaves consistently across conversations	USP (LMSYS-USP)
Population-level diversity	Simulator reproduces full distribution of user types	UserLM/Flipping

6. Deep Dive: HumanLM (Wu et al., Stanford/NYU/Accenture, Feb 2026)

“Human Language Models: Generating Realistic Human Responses in Conversations”

The Core Thesis

HumanLM is the most theoretically grounded paper in the collection. Its central claim: **the right objective for user simulation is state alignment, not response imitation**. Everything else follows from this.

The HUMANUAL Benchmark

Before training, the paper constructs **HUMANUAL** — a benchmark specifically designed to evaluate user simulation along state-aligned dimensions. This is important:

most prior evaluation was ad hoc or borrowed from dialogue generation metrics. HUMANUAL makes the evaluation principled.

Scale: 6 datasets $\times$ $\sim 23,000$ users $\times$ $\sim 227,000$ responses.

Dataset composition: | Dataset | Domain | Characteristics | |----|----|-----
 -| | News | News article discussion | Belief expression, stance | | Book | Book discussion/recommendation | Preference, value expression | | Opinion | Opinionated discussion | Stance, value, argumentation | | Politics | Political discussion | Strong priors, belief update resistance | | Chat | Casual conversation | Emotion, communication style | | Email | Email writing assistance | Goal, communication style |

The domain diversity is deliberate: different domains stress-test different dimensions of the state space. Political discussion is specifically designed to test *belief resistance* — real humans don’t update political beliefs based on one AI response; a simulator that does is unrealistic.

Evaluation dimensions: For each (model, dataset) pair, HUMANUAL evaluates: 1. **Response Quality** (general naturalness) 2. **Belief Consistency** (does the user’s expressed beliefs stay coherent?) 3. **Goal Achievement** (is the user’s conversational goal served?) 4. **Stance Stability** (does stance shift appropriately, not arbitrarily?) 5. **Emotional Appropriateness** (are emotional expressions contextually fitting?) 6. **Style Coherence** (is the communication style consistent?)

GRPO Training

The training objective is **Group Relative Policy Optimization**, adapted from the RL literature:

Setup: - Policy model π : the user simulator. - Reference model π_{ref} : frozen base model (prevents reward hacking). - Reward model r : scores the state-alignment of generated responses. - KL penalty: keeps π close to π_{ref} to prevent degenerate solutions.

GRPO specifics: Rather than single-sample policy gradient, GRPO generates G responses per prompt (a “group”), computes relative rewards within the group, and uses the relative rankings as the gradient signal. This reduces variance compared to standard REINFORCE and avoids the need for a separate critic/value network (unlike PPO).

The reward decomposition: - **State consistency reward:** Given the conversation history and the model’s predicted internal state, does the generated response match the predicted state? - **Cross-turn state coherence reward:** Is the state at turn t consistent with the state at turn $t-1$ (given the conversational events between them)? - **Human distributional reward:** Does the generated response stay within the distribution of actual human responses (measured against HUMANUAL)?

The latent state inference problem: States are never directly observed in training data. HumanLM uses a two-stage approach: 1. **State inference model** (separate): A model trained to infer the six-dimensional state from conversation history, trained on a smaller annotated subset. 2. **Generation policy** (trained via GRPO): Conditions on inferred state.

This creates a dependency: the quality of the generation policy is upper-bounded by the quality of the state inference model. The paper acknowledges this but argues the inference model can be improved independently.

Key Experimental Results

Against prior baselines (GPT-4 prompted, fine-tuned Llama-3 without state alignment): - HumanLM significantly outperforms on **all 6 HUMANUAL evaluation dimensions**. - Biggest gains on **Political** (belief resistance) and **Opinion** (stance stability) — exactly the domains where state dynamics matter most. - Smallest gap on **Chat** (casual conversation) — where style consistency is the dominant factor and state dynamics matter less.

Ablation findings: - Removing state conditioning → large drop on belief/stance dimensions, small drop on style. - Removing GRPO → modest drop across all dimensions (SFT alone is insufficient). - Removing cross-turn coherence reward → largest single component ablation drop.

The cross-turn coherence reward finding is important: **it's the multi-turn dynamics that state alignment buys you, not just better single-turn responses**. This makes intuitive sense — the value of having a coherent internal state is specifically that it persists and evolves across turns.

Critical Assessment

Strengths: - Theoretically principled: the six-dimensional state space is grounded in psychology. - HUMANUAL is a genuine contribution — a benchmark that evaluates what actually matters. - GRPO training framework is technically sound.

Open questions: - State annotation bottleneck: GRPO requires a reward model that evaluates state consistency, which requires annotated states, which requires human annotators who can ground truth what a user's "belief" is at each turn. This is hard to scale. - The six dimensions are a modeling choice, not a universal truth. Political beliefs and spiritual beliefs both fit "value" but are psychologically quite different. - Evaluation on HUMANUAL: the paper shows HumanLM wins on the benchmark it partially helped design. Held-out evaluation on external datasets would strengthen the case. - No downstream RL experiment: the paper argues state alignment matters for RL training, but never actually trains an RL agent against HumanLM to demonstrate improved downstream performance.

7. Deep Dive: Flipping the Dialogue / UserLM (Naous, Laban, Xu, Neville — Microsoft Research, Oct 2025)

"Flipping the Dialogue: Language Models in the User Seat"

Framing Contribution

The most important thing this paper does is **reframe the problem at the field level**, not just propose a specific method. The “flipping” is: instead of asking “how do we get existing LLMs to simulate users better,” ask “what would a purpose-built User LM look like?”

The analogy is to how the field developed specialized models for other non-default tasks (code generation, mathematical reasoning, etc.). Those required purpose-built data, objectives, and evaluation. User simulation deserves the same treatment.

UserLM-8b: Architecture and Training

Base: Llama 3.1-8B, **base checkpoint** (critically: not Llama 3.1-8B-Instruct).

Training data: WildChat — a dataset of ~1 million real human-AI conversations collected from public Gradio deployments of various open-source LLMs. Only **user turns** are used for training. Assistant turns are used only for context (they appear in the input but not in the training targets).

Data preprocessing: - Filter for English conversations. - Remove conversations where the “user” is clearly automated (bot-generated queries). - Deduplicate near-duplicate conversations. - Subsample to preserve topic and style diversity.

Training objective: Standard language modeling (next-token prediction) on user turns, conditioned on preceding assistant turns.

Why base, not instruct? This is the paper’s most empirically grounded contribution. They run a controlled experiment: - UserLM-base: Llama 3.1-8B (base) fine-tuned on WildChat user turns. - UserLM-instruct: Llama 3.1-8B-Instruct fine-tuned on the same WildChat user turns.

Across all 6 intrinsic evaluation dimensions, **UserLM-base > UserLM-instruct**. The gap is not trivial — it’s consistent and statistically significant.

The interpretation: instruction tuning instills a *style* of output generation (formatted, helpful, verbose, ending with “let me know if you need anything”) that persists even after fine-tuning on user-turn data. The base model is more plastic — it can more cleanly learn the user-turn distribution without fighting against baked-in priors.

The 6 Intrinsic Evaluation Dimensions

Unlike HumanLM’s state-space approach, UserLM evaluates through 6 capability-level dimensions:

1. **Perplexity (PPL)** on held-out user turns from WildChat and other datasets. Measures distributional match — does the model assign high probability to things real users say?
2. **Style Diversity** — does the model generate stylistically varied responses across different simulated users? (Measured via vocabulary diversity, sentence length distribution, etc.)

3. **Goal Pursuit** — does the simulated user stay on topic and pursue a coherent conversational goal across turns? (Measured via semantic similarity between user goal description and conversation trajectory.)
4. **Emotional Realism** — are expressed emotions appropriate to the conversational context? (Measured via an emotion classifier applied to generated turns vs. reference turns.)
5. **Belief Expression** — does the model express consistent, coherent beliefs across a conversation? (Measured via NLI-based consistency checking.)
6. **User-Specificity** — given a persona description, does the model generate turns that are distinguishable from turns generated for a different persona? (Measured via classifier accuracy on persona vs. random-persona outputs.)

WildChat as Training Data: Implications

WildChat is an unusual training corpus in an important way: it captures *naturalistic* human LLM use, including: - Lazy, shorthand queries (“tldr this”) - Frustrated follow-ups (“that’s not what I asked”) - Off-task tangents - Creative and unconventional requests - Non-English and code-switched conversations (filtered out here but present in raw data)

This is different from curated conversational datasets, which tend to overrepresent “cooperative user” behaviors. The WildChat-trained UserLM thus inherits realistic user behaviors that don’t appear in cleaner datasets.

Limitation: WildChat reflects users of *open-source LLMs* — a specific, tech-savvy, self-selected population. The user distribution may not generalize to the full spectrum of real-world LLM users (non-technical users, non-English speakers, enterprise users, etc.).

The Population-Level Contribution

A key empirical contribution of UserLM is demonstrating that it preserves **population-level diversity**. Rather than collapsing to an “average user,” UserLM generates different behaviors for different sampled personas.

The paper measures this via **distributional coverage**: what fraction of the real WildChat user-turn distribution falls within the high-probability region of the simulator’s distribution? UserLM significantly outperforms prompted GPT-4, which tends to oversample from a narrow “helpful cooperative user” mode.

Critical Assessment

Strengths: - Clean experimental design (base vs. instruct ablation is compelling). - WildChat is a good choice for naturalistic user diversity. - Population-level diversity measurement is a genuine methodological contribution. - UserLM-8b is publicly released — concrete research artifact.

Open questions: - No state-alignment: UserLM learns to *imitate* user turns (albeit better than prior work). HumanLM’s critique applies: does it have a coherent internal state that persists across turns, or is each turn generated independently conditioned on history? - WildChat bias: Tech-savvy user population. How does performance degrade when simulating non-technical users? - 8B parameters: Is this a fundamental size? How does UserLM scale? Would a larger model simply over-fit the WildChat distribution more, or would it genuinely improve? - Evaluation breadth: The 6 dimensions are a reasonable choice but not universal. The paper would benefit from downstream evaluation (e.g., train an RL agent against UserLM and measure transfer to real users).

8. Deep Dive: Know You First and Be You Better / USP (Wang et al., CUHK/NUS, ACL 2025)

“Know You First and Be You Better: Modeling Human Profiles for Personalized User Simulation”

The Problem This Addresses

Both HumanLM and UserLM focus on general user simulation — how to reproduce the *distribution* of user behaviors. USP addresses a different sub-problem: **personalized user simulation** — how to simulate a *specific* user given information about that user.

This matters for: - **Personalization research:** You want to evaluate whether a new personalization algorithm actually adapts to individual users. You need a simulator that can be “User A” (a cautious, risk-averse person who asks lots of clarifying questions) vs. “User B” (a fast-moving power user who tolerates some ambiguity). - **Red-teaming specific user types:** Simulate a user with specific demographic characteristics, knowledge level, or vulnerability. - **Data augmentation for specific user cohorts:** Generate training data for rare user types.

The Profile Schema: OF + SC

The core architectural choice is an **implicit profile schema** with two components:

Objective Facts (OF): Stable, verifiable facts about the user: - Demographic information (age range, occupation, education level) - Domain expertise (“knows Python but not Rust”) - Usage patterns (“uses the AI mainly for writing assistance”) - Factual preferences (“prefers metric units”)

Subjective Characteristics (SC): Inferred psychological traits: - Communication style (“tends to be indirect about what they actually want”) - Emotional patterns (“expresses frustration through sarcasm”) - Value priorities (“strongly values privacy over convenience”) - Cognitive patterns (“tends to over-explain context before asking the real question”)

This OF + SC decomposition is important for two reasons: 1. **Extraction tractability:** OF is easier to extract reliably; SC requires more nuanced inference. 2. **Gener-**

alization: OF facts can often be obtained from user profiles/metadata; SC requires behavioral signal.

Profile extraction pipeline: - Input: A collection of conversations involving a specific user (from LMSYS-Chatbot Arena or similar). - Extractor model: A prompted LLM (GPT-4) that reads the conversations and outputs structured OF + SC. - Output: A natural language profile paragraph, plus structured fields.

The extraction quality is evaluated by asking: given the extracted profile, can a reader predict how the user would respond to a new query? (Evaluated via human annotators.)

LMSYS-USP Dataset

The paper constructs **LMSYS-USP** from LMSYS-Chatbot Arena data: - **94,874 conversations** across users who participated in multiple conversations. - For each user with $\geq N$ conversations, hold out some conversations for evaluation and use the rest for profile extraction. - Final dataset: profiles + associated conversation collections for training and evaluation.

Why LMSYS-Chatbot Arena? It’s one of the few public datasets where: 1. Users are identifiable across multiple conversations (via session IDs). 2. Conversations are naturalistic (users exploring real questions, not crowd-sourced tasks). 3. Scale is sufficient for learning cross-session user identity.

Training: Conditional SFT + RLCC

Stage 1: Conditional SFT Standard fine-tuning of a base LLM on user turns, but conditioned on the extracted profile (OF + SC). The model learns: given “this is the kind of person who [SC description]” and “here are their facts [OF description]”, generate responses that fit this profile.

This is straightforward but has a key limitation: the SFT objective has no guarantee that the *generated* utterances would actually lead an outside observer to infer the given profile. The model might satisfy local consistency (each turn sounds like the person) without satisfying global consistency (a conversation-level observer would reconstruct the profile).

Stage 2: RLCC (RL with Cycle Consistency)

The key innovation. The training loop: 1. **Profile → Dialogue generation:** Given profile P , the simulator generates a synthetic dialogue D . 2. **Dialogue → Profile extraction:** The profile extractor (the same GPT-4-based extractor used to build LMSYS-USP) extracts a profile P' from D . 3. **Cycle consistency reward:** $r = \text{sim}(P, P')$ — how similar is the extracted P' to the original P ? (Measured via embedding cosine similarity or ROUGE-based overlap.) 4. **Policy update:** Reinforce the simulator to maximize r .

The cycle consistency objective forces the simulator to generate conversations that are **evidentially sufficient** for the profile. A simulator that generates neutral, non-

committal utterances might look fine locally but would get a low $P' \rightarrow$ low reward, because the profile extractor can't infer the profile from a bland conversation.

Diverse Profile Sampling (GKDE + UMAP)

A nice methodological contribution for training data construction: using GKDE (Gaussian Kernel Density Estimation) in UMAP-projected profile space to ensure training profiles are diverse. Without this, profiles in training data cluster around common user types (the “helpful, curious person with no strong opinions” attractor). GKDE sampling over-samples rare profile regions to improve the model's ability to simulate edge-case users.

DPC: Dialogue Profile Consistency Metric

The paper introduces **Dialogue Profile Consistency (DPC)** as an evaluation metric: 1. Given a generated conversation, use FactScore-style atomic fact decomposition to extract specific factual and behavioral claims. 2. Check each claim against the ground-truth profile. 3. DPC = fraction of extracted claims that are supported by the profile.

This metric is more fine-grained than human evaluation and more reliable than embedding similarity. It's also composable: you can compute DPC(OF) and DPC(SC) separately to understand where the model succeeds and fails.

Key finding: DPC(SC) is harder than DPC(OF) — simulating subjective characteristics (communication style, emotional patterns) is harder than simulating objective facts. This makes intuitive sense but is good to have quantified.

Critical Assessment

Strengths: - OF + SC profile schema is well-motivated and practically useful. - RLCC is a clever training objective that addresses the global consistency problem. - LMSYS-USP is a genuine dataset contribution. - DPC metric is a solid contribution to the evaluation toolkit.

Open questions: - **Profile extraction quality is a bottleneck.** The whole pipeline depends on GPT-4 being able to accurately extract OF + SC. For unusual users or sparse conversation histories, this may fail. The quality of the extracted profile is never perfect, which means the reward signal in RLCC is noisy. - **Privacy concerns.** The LMSYS dataset involves real user conversations. Building user profiles from these is potentially sensitive. The paper doesn't address privacy. - **Cross-session generalization.** LMSYS-USP is built on LMSYS-Chatbot Arena, where “session” is loosely defined. Can USP simulate a user across truly long-term interactions (months of usage history)? - **Profile drift.** Real users change over time. USP assumes a static profile. How should the profile be updated? - **The extractor loop.** RLCC requires the same extractor to be used for both training (constructing reward signal) and evaluation (evaluating DPC). This risks teaching the simulator to fool the *specific extractor*, not to genuinely exhibit profile-consistent behavior.

9. Deep Dive: ChatBench (Chang, Anderson, Hofman — Microsoft Research, ACL 2025)

“ChatBench: A User Simulator Benchmark for Human-AI Interaction Evaluation”

The Motivation: Static Benchmarks Are Wrong

ChatBench starts from a simple but powerful observation: **the accuracy of an AI system measured on static benchmarks (MMLU, TriviaQA, etc.) poorly predicts how well that system does when a human is asking the questions.**

Why? Because static benchmarks assume perfect questions — well-formed, unambiguous, with clear correct answers. Real users don’t ask perfect questions. They ask:

- Ambiguous questions that require clarification.
- Questions with implicit assumptions that need surfacing.
- Questions with wrong premises that the AI needs to gently correct.
- Sequences of related questions where later questions depend on understanding earlier answers.

The core empirical finding: AI systems that perform identically on static benchmarks can differ enormously in user-AI collaborative accuracy — the accuracy achieved when a simulated user interacts with the AI in multi-turn conversation to solve the same underlying problem.

The MMLU → Conversation Conversion

ChatBench converts static MMLU items into user-AI conversations by:

1. Given an MMLU question and its answer, design a “user” who wants to answer this question but doesn’t know the answer.
2. The user simulator conducts a conversation with the AI to extract the answer.
3. Measure whether the user correctly identifies the right answer at the end of the conversation.

The conversion is non-trivial: the user must be realistic (not simply parroting the MMLU question), and the conversation must preserve the difficulty of the original question (not trivially simplify it to a single direct request).

User simulator used: A fine-tuned user simulator (not specified in detail in the 10 pages I read), trained on real human-AI conversations from comparable question-answering settings.

Key Findings

Finding 1: AI-alone accuracy (zero-shot answering the MMLU question directly) has only weak correlation with user-AI accuracy. The R^2 is surprisingly low — less than 0.3 in several settings.

Finding 2: Adding a fine-tuned user simulator improves correlation significantly — 22-26 percentage points in R^2 depending on the domain. This validates the approach: a good user simulator makes benchmark scores more predictive of real-world performance.

Finding 3: Some AI systems have high static accuracy but low collaborative accuracy. These tend to be systems that give technically correct but poorly explanatory answers — an expert-user interaction problem. Conversely, some systems with lower static accuracy achieve high collaborative accuracy because they’re better at understanding what the user is actually trying to do.

Finding 4: Domain matters enormously. User-AI collaborative accuracy is much higher in domains where the AI can explain its reasoning incrementally (math, science) vs. domains requiring specialized knowledge retrieval (history, geography).

Why This Matters for the Field

ChatBench is the clearest demonstration to date that **evaluation infrastructure is broken** and that user simulators are the fix. The paper makes the case that:

1. Building better AI systems requires better evaluation.
2. Better evaluation requires realistic user simulation.
3. Therefore, user simulation is upstream of AI system progress.

This is a different argument than HumanLM’s (training requires state-aligned simulators) or UserLM’s (a new model class is needed). It’s an *evaluation* argument: even if you don’t care about training, you should care about user simulation because your benchmarks are misleading you.

Critical Assessment

Strengths: - The core finding (static benchmarks poorly predict collaborative accuracy) is important and clearly demonstrated. - MMLU → conversation conversion is a principled methodology. - The 22-26 point correlation improvement from fine-tuned simulators is a concrete result.

Open questions: - The paper is about *benchmarking*, not about *building* user simulators. It motivates the need but doesn’t provide a blueprint. - MMLU is a multiple-choice benchmark with objectively correct answers. Generalizing to open-ended domains (creative writing assistance, emotional support) is unclear. - The fine-tuned simulator used is not fully described. How sensitive are the results to simulator quality? - The paper’s user simulator is domain-specific (question-answering). How much of the evaluation framework generalizes?

10. Jessie Lin’s Framework: What Does It Take?

“User simulators bridge RL with real-world interaction” (Blog Part I) “What does it take to build a human-like user simulator?” (Blog Part II) Jessie Lin & Nick Tomlin, Jul-Sep 2025

These two blog posts are the most accessible high-level framing of the field and serve as a useful conceptual compass. They’re worth reading as a pair.

Part I: Five Failure Modes (the diagnostic)

Lin identifies five specific ways current LM-based user simulators fail. I’ve described these under Generation 3 (§2), but they deserve reinforcement:

1. Too knowledgeable: The simulator “knows” things its persona wouldn’t know. This isn’t just about factual knowledge — it’s about *meta-knowledge* (knowing what the AI knows, understanding how AI systems work) that a naive user wouldn’t have. An LLM playing a “confused elderly user” still knows the right terminology for what it’s confused about.

2. Perfect memory: Real humans have cognitive load limits. They misremember what was said earlier, ask about things that were already explained, and don’t maintain perfect context across long conversations. Simulators based on long-context LLMs have essentially perfect recall.

3. Too helpful: This is the alignment leakage problem. Instruction-tuned LLMs cooperate with the AI agent even when their persona would resist, argue back, or refuse. The helpfulness bias is deep and hard to suppress via prompting.

4. Inconsistent beliefs: Without a persistent state representation, the same simulator may express contradictory views in different turns or different sessions. This is especially visible in longer conversations.

5. Shallow acting: The most important failure mode. The simulator produces the *surface behavior* of a user type without having the underlying generative model. It’s like an actor who knows all the lines but doesn’t understand the character’s motivation — fine for short scenes, breaks down in improvisation or novel situations.

The RL environment design framing: Lin argues that these failure modes aren’t just quality problems — they *misshape* the RL training signal. An RL agent trained against a too-helpful, perfect-memory, too-knowledgeable simulator will learn exploits that work against the simulator but not against real users. This is the same phenomenon as reward hacking in RLHF — the model learns to satisfy the proxy reward rather than the true objective.

Part II: Context + Scaffold + Objective (the constructive framework)

Having diagnosed the failures, Lin and Tomlin propose a framework for building better simulators. Three components:

Context: What information does the simulator have access to? - *Persona/profile*: Background of the simulated user (demographics, expertise, goals). - *Task context*: What is the user trying to accomplish in this conversation? - *Prior interaction history*: What has happened in previous sessions with this user? - *World knowledge*: What does this user know about the topic (different from what the LLM knows)?

The challenge is calibrating context appropriately — the simulator should know what *its character* knows, not what the LLM knows. This requires either careful prompting or architectural separation.

Scaffold: What computational structure does the simulator use to generate coherent behavior?

This is where Lin and Tomlin connect to the state alignment literature: - *Explicit goal tracking*: The simulator maintains an explicit representation of what the user is trying to achieve and checks each generated turn against it. - *Belief state maintenance*: The simulator has a belief state that updates as the AI provides information. - *Self-knowledge limits*: The simulator has an explicit model of what it doesn't know (and sometimes doesn't know that it doesn't know). - *Memory & cognitive load*: The simulator has bounded memory — earlier information degrades, gets confused with similar information. - *Influence/belief update modeling*: How should the user's beliefs update in response to AI persuasion? (Bayesian updating? Motivated reasoning? Anchoring bias?)

Objective: What is the simulator optimized to produce? - *Behavioral realism*: Does it look like a real human? - *Goal fidelity*: Does it pursue its stated goal effectively? - *Downstream task performance*: Does an RL agent trained against this simulator perform better with real humans?

The key insight: **these three objectives are in tension**. Behavioral realism might require generating confused, off-task turns; goal fidelity requires staying on-task. The right balance depends on what the simulator is *for*.

Lin's Open Questions

The second blog post ends with a set of open research questions that are worth cataloging as a research agenda:

1. How do we calibrate the simulator's knowledge to match its persona's actual knowledge state? (Without expensive human annotation of every persona's knowledge profile.)
2. How do we implement realistic memory degradation without making the simulator unusably inconsistent?
3. Is there a unified framework that handles both task-oriented and open-domain conversation? (The former needs explicit goal tracking; the latter needs style/personality consistency.)
4. How do you evaluate a simulator other than by running actual RL training? (The gold standard evaluation is: does an RL agent trained against this simulator improve against real users? This is expensive.)
5. Can simulators themselves learn from interaction — i.e., can a user simulator update its internal model of “what a particular user is like” from ongoing interaction, the way a real user updates their communication strategy over time?

11. Supporting Papers: Quick Takes

Beyond the six core papers, the UserLM.Systems collection includes additional papers that are worth mapping:

IQA-EVAL

Interactive Question Answering Evaluation - Focus: evaluation of QA systems via simulated interaction. - Key contribution: Shows that simulated multi-turn QA reveals failure modes invisible to static evaluation — similar thesis to ChatBench but predating it. - Relationship to core papers: Forerunner of ChatBench; different methodology (no MMLU conversion).

DuetSim

Dual-Role Simulation in Dialogue - Focus: Simulating both sides of a dialogue (user and system) simultaneously. - Key innovation: Trains a single model to play both roles, enabling faster data generation. - Limitation: The two roles are not independent, which creates potential for mode collapse.

PlatoLM

Plato Language Model - Focus: Training on philosophical dialogue to model users who engage in Socratic-style questioning. - Niche but interesting: Demonstrates that conversational *mode* (not just topic) can be simulated.

Reliable LLM-based User Simulator

- Focus: Reducing simulator inconsistency in task-oriented dialogue.
- Key technique: Explicit dialogue state tracking to enforce consistency.
- Relationship to core papers: More closely aligned with HumanLM’s state tracking goal, but applied to task-oriented rather than open-domain conversation.

Multi-trait mTAD

- Focus: Multi-trait user simulation for task-aware dialogue.
- Key contribution: Demonstrates that users have multiple simultaneous traits (personality, expertise, communication style) that need to be jointly modeled, not treated independently.
- Relevant to USP: Multi-trait modeling vs. USP’s OF+SC decomposition.

Parrot

- Focus: Automatically generating diverse user personas for dialogue evaluation.
- Contribution: A pipeline for persona generation, not simulation per se.
- Useful as upstream: Could supply profiles to USP or HumanLM.

In-context Learning Simulators

- Multiple papers on using few-shot examples to adapt LLMs to user simulation.
- Represent the Gen-3 approach, now largely superseded by Gen-4 fine-tuned models.

Domain-Independent Simulators

- Focus: Building simulators that generalize across conversational domains without domain-specific fine-tuning.
- Key challenge: Domain-specific vocabulary and interaction patterns make generalization hard.
- Relationship to UserLM: UserLM’s WildChat training covers many domains; this is the right empirical approach.

Unified Dialogue Simulator

- Focus: A single model that can simulate users across task-oriented and open-domain dialogue.
 - Key contribution: Architecture for unified simulation; trade-offs between specialization and generalization.
-

12. Open Problems and Research Frontiers

Problem 1: The Evaluation Bootstrapping Problem

The problem: The gold-standard evaluation of a user simulator is *downstream*: does an RL agent trained against this simulator perform better with real humans? But running this evaluation is expensive (requires deploying real RL agents and measuring real user satisfaction). All the evaluation metrics in current papers (PPL, DPC, human ratings, LM-judge) are *proxies* for this real evaluation.

Why it matters: Proxy metrics can be gamed. A simulator can achieve high DPC (profile consistency) by generating utterances that explicitly mention profile attributes, which a human would never do. A simulator can achieve low PPL by generating high-frequency user patterns (short queries, simple follow-ups) without ever generating rare but important user behaviors.

Research directions: - Develop cheaper approximations of the downstream evaluation (e.g., use human evaluators to rate conversation quality in a structured way that’s shown to predict downstream RL performance). - Build shared evaluation infrastructure so different papers’ simulators can be compared on the same downstream tasks. - Explore whether there are *leading indicators* — cheap metrics that reliably predict downstream performance — analogous to perplexity as a leading indicator for language model performance.

Problem 2: Knowledge Calibration

The problem: The simulator should know what its persona knows — not what the training LLM knows. But the training LLM was trained on vast amounts of human knowledge, which is hard to suppress.

Why it matters: A user simulator that “knows too much” gives misleading feedback to RL agents. An RL agent trained against an omniscient simulated user will not

learn to handle the gaps and misunderstandings that characterize interactions with real users who don't know everything.

Research directions: - **Selective forgetting/unlearning:** Techniques to suppress specific knowledge from a fine-tuned model. - **Knowledge profile conditioning:** Explicitly represent what the simulated user knows as a structured knowledge profile, and condition the simulator on this profile. - **Retrieval-augmented simulation:** The simulator only has access to knowledge explicitly retrieved from a user-specific knowledge base, rather than relying on the LLM's parametric knowledge.

Problem 3: Bounded Memory and Cognitive Load

The problem: Real users have limited working memory, misremember conversation history, and lose context over long conversations. LLM-based simulators have essentially perfect recall within the context window.

Why it matters: RL agents trained against perfect-memory simulators will not learn to: - Repeat important information proactively. - Check for understanding. - Handle users who refer back to earlier parts of the conversation imprecisely.

Research directions: - **Explicit memory degradation models:** Implement a structured model of memory decay (e.g., information decays as a function of time and interference from similar information). - **Cognitive load modeling:** Track the total cognitive burden of the conversation and adjust coherence/memory accordingly. - **Uncertainty simulation:** The simulator should sometimes "not remember" or "be unsure" about what was said earlier — generating uncertain, approximate references to earlier conversation.

Problem 4: Belief Dynamics and Persuasion

The problem: How does a simulated user update their beliefs in response to AI-provided information? Real users have a complex mix of rational updating, anchoring bias, confirmation bias, and motivated reasoning. Simulators either don't update (stubborn) or update too readily (sycophantic).

Why it matters: For RL training in domains involving advice-giving (medical, financial, legal), the user simulator's belief update model determines whether the RL agent learns to explain clearly vs. just assert confidently.

Research directions: - **Bayesian belief update models:** Explicitly model prior beliefs and likelihood ratios; update beliefs probabilistically. - **Bias simulation:** Train simulators to reproduce human cognitive biases (anchoring, availability heuristic, confirmation bias) at the population level. - **Persuasion susceptibility modeling:** Model individual users' resistance to persuasion as a function of their profile (people with high dogmatism update beliefs less readily).

Problem 5: Emotional Realism and Dynamic Emotional Arcs

The problem: Current simulators model emotion as a snapshot — the emotion *at* each turn. Real conversations have emotional arcs: frustration builds over failed

attempts, trust develops over time, satisfaction or disappointment at the resolution.

Why it matters: RL agents that learn only from snapshot emotions will learn local recovery (fix the frustration in the current turn) but not longitudinal emotional management (prevent frustration from building in the first place).

Research directions: - **Emotion trajectory models:** Model emotion as a dynamical system with state and transition dynamics, not just a per-turn label. - **Expectation violation models:** Model how violations of user expectations (the AI did something surprising, good or bad) affect the emotional arc. - **Recovery modeling:** Simulate how users respond to AI attempts to recover from errors.

Problem 6: Cross-Session Identity Continuity

The problem: Real users are the same person across multiple sessions. They remember (imperfectly) what the AI told them before, they build a relationship with the AI system over time, and their communication style with the AI adapts based on experience.

Why it matters: AI systems designed for long-term engagement (productivity tools, health coaches, tutors) need simulators that can model multi-session dynamics.

Research directions: - **Episodic memory:** The simulator maintains an episodic record of previous sessions, from which it samples (with realistic forgetting) to inform current session behavior. - **Relationship state:** A dimension of the user state that represents the user's relationship with the AI (trust, familiarity, communication norms). - **Style adaptation:** Model how users' communication style changes as they become more experienced with the AI system.

Problem 7: Adversarial and Edge-Case Users

The problem: Current simulators tend to model cooperative, goal-directed users. Real users include: - Users with conflicting or under-specified goals. - Users who test the AI's limits. - Users who are malicious (jailbreaking, prompt injection). - Users with mental health challenges or under acute distress. - Non-native speakers with unconventional phrasings.

Why it matters: AI safety and robustness depend on training against these user types. Current simulators don't cover them.

Research directions: - **Adversarial simulation:** Specialized simulators trained on jailbreak attempts and limit-testing behaviors. - **Vulnerability-aware simulation:** Simulators that model users in distress, with appropriate emotional and cognitive profiles. - **Linguistic diversity:** Simulators trained on non-native speaker data, code-switched language, unconventional dialects.

13. Synthesis: Where the Field Is Heading

The Three Bets

The field is currently organized around three distinct research bets, embodied in the three main technical papers:

Bet 1 (HumanLM): State alignment is the key. The right framing is psychological: build a model of the user’s internal state (beliefs, goals, values, stance, emotion, style) and generate responses from that state. This requires new training objectives (GRPO), new evaluation frameworks (HUMANUAL), and new architectural choices (state inference + generation). The payoff is a simulator with the right *structure* for long-horizon reasoning about user behavior.

Bet 2 (UserLM/Flipping): Purpose-built User LMs are the key. The right framing is model identity: stop fighting against assistant-aligned LLMs and build a language model whose fundamental identity is the user side of conversations. This requires using base model checkpoints, training on user-turn data, and evaluating on user-specific capabilities. The payoff is a reusable model class that the ecosystem can build on.

Bet 3 (USP/Know You First): Personalization is the key. The right framing is individuation: the goal isn’t to simulate a generic user but to simulate *this specific user*. This requires profile extraction (OF + SC), profile-conditioned generation, and cycle-consistent training (RLCC). The payoff is simulators that can model real individuals rather than demographic archetypes.

These bets are **not mutually exclusive**. The ideal system would be all three: a User LM (Bet 2), trained with state alignment (Bet 1), that can be conditioned on individual profiles (Bet 3). The three research threads are converging; the next generation of papers will likely integrate them.

The Integration Hypothesis

The theoretical integration looks like this:

```

Profile (OF + SC)
  ↓
State Initialization
(belief, goal, value, stance, emotion, style)
  ↓ ← base User LM
State-Conditioned Generation
  ↓
Simulated User Turn
  ↓ ← RL
State Update (belief update from AI response, etc.)
  ↓
Next Turn...
```

The state space provides *continuity* (HumanLM’s contribution); the profile provides *individuation* (USP’s contribution); the base User LM provides the right *prior* for

user-like generation (UserLM’s contribution).

The missing piece: GRPO and RLCC are training methods for specific objectives. What’s missing is a *unified training objective* that simultaneously optimizes: - Distributional realism (PPL on real user turns) - Profile consistency (DPC from USP) - State coherence (HumanLM’s cross-turn reward) - Downstream RL performance (the real ground truth)

This four-way trade-off is the next fundamental research challenge.

The Ecosystem Bet

Beyond the technical integration, there’s a meta-level bet about the *ecosystem*. UserLM explicitly argues for a new model class. If this bet is right, we’ll see:

1. **Foundation User LMs:** Large-scale base user simulators (like GPT-4 or Llama, but for the user side), which practitioners fine-tune for specific domains.
2. **User LM benchmarks:** Standardized evaluation suites (HUMANUAL is one start) that become the “MMLU” of user simulation.
3. **User LM APIs:** Services analogous to embedding APIs or completion APIs, specifically for simulating user conversations.
4. **RL training infrastructure:** Toolkits that make it easy to train dialogue agents against User LMs, with built-in population diversity and behavioral realism.

The ChatBench paper provides the strongest argument for why this ecosystem matters: better evaluation infrastructure leads to better AI systems, and better evaluation requires realistic user simulation. If ChatBench’s methodology (static benchmarks don’t predict collaborative accuracy; user simulators close the gap) becomes the field consensus, there will be strong incentive to build and maintain these User LM foundations.

The Alignment Implications

There’s a deeper issue lurking in this literature that no single paper fully addresses: **the alignment implications of user simulators.**

If we use user simulators to train AI systems, the AI’s behavior will be optimized against the simulator’s behavior. If the simulator models a specific population of users (WildChat users, LMSYS-Chatbot Arena users), the AI will become good at serving *that population* and potentially less good at serving others.

The diversity framing in UserLM (GKDE sampling, population coverage) is a step toward addressing this, but it’s incomplete. True population coverage would require simulators trained on the full diversity of global internet users, including: - Non-English speakers (most of the world’s internet users). - Users with low digital literacy. - Users in low-resource settings. - Older and younger users with different interaction patterns.

Building user simulators from LMSYS-Chatbot Arena and WildChat systematically under-represents these populations. The user simulators we build will reflect the

users who used early LLM chatbots — who are not a representative sample of the people AI systems will serve.

This is the alignment problem for user simulation: **misspecification of the user population leads to misspecification of the training signal, which leads to AI systems optimized for the wrong people.**

Appendix: Quick Reference Tables

Papers in UserLM.Systems (Full Collection)

#	Title (Short)	Authors	Year	Focus	Key Contribution
1	HumanLM	Wu et al.	2026	State-aligned simulation	HUMANUAL benchmark, GRPO training, 6-dim state space
2	Flipping the Dialogue	Naous et al. (MSR)	2025	User LM as a model class	UserLM-8b, base > instruct result, WildChat training
3	Know You First	Wang et al.	2025	Personalized simulation	USP framework, OF+SC profiles, RLCC, DPC metric
4	ChatBench	Chang et al. (MSR)	2025	Evaluation	Static benchmarks $\neq$ collaborative accuracy; +22-26pt R^2
5	User Simulators Bridge RL	Lin & Tomlin	2025	Framework/evaluation	5 failure modes; RL environment design problem
6	What Does It Take	Lin & Tomlin	2025	Framework/evaluation	Construct Scaffold + Objective framework
7	IQA-EVAL	—	—	Evaluation	Interactive QA evaluation
8	DuetSim	—	—	Dual-role simulation	Simulating both user and system
9	PlatoLM	—	—	Specialized simulation	Socratic questioning mode
10	Reliable LLM User Sim	—	—	Task-oriented	Explicit state tracking for consistency
11	Multi-trait mTAD	—	—	Multi-trait	Joint modeling of multiple user traits
12	Parrot	—	—	Persona generation	Automated diverse persona generation

#	Title (Short)	Authors	Year	Focus	Key Contribution
13+	In-context learning, domain-independent, unified	Various	Various	Various	Gen-3 approaches + generalization

Key Metrics at a Glance

Metric	Paper	What It Measures	How Computed
PPL	UserLM/Flipping	Distributional match to real user turns	LM perplexity on held-out human turns
DPC	USP	Profile consistency of generated dialogue	FactScore atomic fact checking vs. profile
HUMANUAL-6	HumanLM	Six-dimensional state alignment	Human eval + LM judge on 6 dimensions
R ² (static→collab)	ChatBench	Benchmark predictiveness gap	Correlation between static acc and user-AI acc
Style Diversity	UserLM	Population coverage	Vocabulary/style distribution breadth

Training Methods at a Glance

Method	Paper	What It Does
SFT on user turns (base)	UserLM/Flipping	Standard fine-tuning with user-turn supervision
GRPO	HumanLM	RL with group-relative rewards for state alignment
Conditional SFT	USP	Profile-conditioned fine-tuning
RLCC	USP	RL with cycle-consistent profile reward

Document compiled from full PDF reads of HumanLM, Flipping the Dialogue, Know You First, ChatBench, and Jessy Lin’s two blog posts, plus metadata from 13+ additional papers in the UserLM.Systems collection.

Avishek’s Zotero → NLP → LLMs → UserLM → Systems